

Install Package

```
%pip install datasets
```

Show hidden output

Import Libraries

```
from datasets import load_dataset, DatasetDict
from transformers import AutoTokenizer, TFAutoModelForSequenceClassification
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

Load Dataset

```
import pandas as pd

pandas_df = pd.read_parquet("hf://datasets/tdavidson/hate_speech_offensive/data/train-00000-of-00001.parquet")
```

pandas_df.head()

	count	hate_speech_count	offensive_language_count	neither_count	class	tweet
0	3	0	0	3	2	!!! RT @mayasolovely: As a woman you shouldn't...
1	3	0	3	0	1	!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2	3	0	3	0	1	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3	3	0	2	1	1	!!!!!!! RT @C_G_Anderson: @viva_based she lo...
4	6	0	6	0	1	!!!!!!!!!!!! RT @ShenikaRoberts: The shit you...

Next steps:

Generate code with pandas_df

View recommended plots

New interactive sheet


pandas_df['class']

	class
0	2
1	1
2	1
3	1
4	1
...	...
24778	1
24779	2
24780	1
24781	1
24782	2

24783 rows × 1 columns

dtype: int64

```
import seaborn as sns
sns.countplot(x='class', data=pandas_df, hue='class')
```



Clean Data

```
pandas_df['tweet_cleaned'] = pandas_df['tweet'].str.replace('@[A-Za-z0-9]+\s?', '', regex=True)
pandas_df.head()
```

	count	hate_speech_count	offensive_language_count	neither_count	class	tweet	tweet_cleaned
0	3	0	0	3	2	!!! RT @mayasolovely: As a woman you shouldn't...	!!! RT : As a woman you shouldn't complain abo...
1	3	0	3	0	1	!!!! RT @mleew17: boy dats cold...tyga dwn ba...	!!!! RT : boy dats cold...tyga dwn bad for cu...
2	3	0	3	0	1	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...	!!!!!! RT Dawg!!!! RT : You ever fuck a bitch...
3	3	0	2	1	1	!!!!!! RT @C_G_Anderson: @viva_based she lo...	!!!!!! RT _G_Anderson: _based she look like...
4	6	0	6	0	1	!!!!!! RT @ShenikaRoberts: The shit you...	!!!!!! RT : The shit you hear about me ...

Next steps: [Generate code with pandas_df](#) [View recommended plots](#) [New interactive sheet](#)

```
from datasets import Dataset
```

```
ds = Dataset.from_pandas(pandas_df)
ds
```

```
Dataset({
  features: ['count', 'hate_speech_count', 'offensive_language_count', 'neither_count', 'class', 'tweet', 'tweet_cleaned'],
  num_rows: 24783
})
```

```
train_test_valid = ds.train_test_split()
```

```
test_valid = train_test_valid['test'].train_test_split()
```

```
train_test_valid_dataset = DatasetDict({
  'train': train_test_valid['train'],
  'test': test_valid['test'],
  'valid': test_valid['train']
})
```

```
dataset = train_test_valid_dataset.remove_columns(['hate_speech_count', 'offensive_language_count', 'neither_count', 'count'])
dataset
```

```
DatasetDict({
  train: Dataset({
    features: ['class', 'tweet', 'tweet_cleaned'],
    num_rows: 18587
  })
  test: Dataset({
    features: ['class', 'tweet', 'tweet_cleaned'],
    num_rows: 1549
  })
  valid: Dataset({
    features: ['class', 'tweet', 'tweet_cleaned'],
    num_rows: 4647
  })
})
```

✓ Tokenizer

✓ load tokenizer

```
tokenizer = AutoTokenizer.from_pretrained("bert-base-cased")
```

```
➦ /usr/local/lib/python3.10/dist-packages/transformers/tokenization_utils_base.py:1601: FutureWarning: `clean_up_tokenization_spaces` was not se  
warnings.warn(
```

✓ Test tokenizer

```
text = "Just checking tokenization"
```

```
output = tokenizer(text)
```

```
output
```

```
➦ {'input_ids': [101, 2066, 9444, 22559, 2734, 102], 'token_type_ids': [0, 0, 0, 0, 0, 0], 'attention_mask': [1, 1, 1, 1, 1, 1]}
```

```
tokens = tokenizer.convert_ids_to_tokens(output['input_ids'])  
tokens
```

```
➦ ['[CLS]', 'Just', 'checking', 'token', '##ization', '[SEP]']
```

```
print(f"Tokenized text: {tokenizer.convert_tokens_to_string(tokens)}")
```

```
➦ Tokenized text: [CLS] Just checking tokenization [SEP]
```

```
print(f"Vocab size is : {tokenizer.vocab_size}")
```

```
print(f"Model max length is : {tokenizer.model_max_length}")
```

```
print(f"Model input names are: {tokenizer.model_input_names}")
```

```
➦ Vocab size is : 28996  
Model max length is : 512  
Model input names are: ['input_ids', 'token_type_ids', 'attention_mask']
```

✓ Tokenize Function

```
def tokenize_function(train_dataset):  
    return tokenizer(train_dataset['tweet_cleaned'], padding='max_length', truncation=True)
```

```
tokenized_dataset = dataset.map(tokenize_function, batched=True)
```

```
tokenized_dataset
```

```
➦ Map: 100%                               18587/18587 [00:14<00:00, 2023.54 examples/s]  
Map: 100%                               1549/1549 [00:01<00:00, 1143.47 examples/s]  
Map: 100%                               4647/4647 [00:02<00:00, 1565.90 examples/s]  
DatasetDict({  
  train: Dataset({  
    features: ['class', 'tweet', 'tweet_cleaned', 'input_ids', 'token_type_ids', 'attention_mask'],  
    num_rows: 18587  
  })  
  test: Dataset({  
    features: ['class', 'tweet', 'tweet_cleaned', 'input_ids', 'token_type_ids', 'attention_mask'],  
    num_rows: 1549  
  })  
  valid: Dataset({  
    features: ['class', 'tweet', 'tweet_cleaned', 'input_ids', 'token_type_ids', 'attention_mask'],  
    num_rows: 4647  
  })  
})
```

✓ Tokenize Dataset

```
train_dataset = tokenized_dataset['train']  
eval_dataset = tokenized_dataset['valid']  
test_dataset = tokenized_dataset['test']
```

```
train_dataset
```

```
Dataset({
  features: ['class', 'tweet', 'tweet_cleaned', 'input_ids', 'token_type_ids', 'attention_mask'],
  num_rows: 18587
})
```

```
train_set = train_dataset.remove_columns(['tweet', "tweet_cleaned"]).with_format('tensorflow')
```

```
tf_eval_dataset = eval_dataset.remove_columns(['tweet', "tweet_cleaned"]).with_format('tensorflow')
```

```
tf_test_dataset = test_dataset.remove_columns(['tweet', "tweet_cleaned"]).with_format('tensorflow')
```

```
train_features = { x: train_set[x] for x in tokenizer.model_input_names }
```

```
train_set_for_final_model = tf.data.Dataset.from_tensor_slices((train_features, train_set['class'] ))
```

```
train_set_for_final_model = train_set_for_final_model.shuffle(len(train_set)).batch(8)
```

```
eval_features = {x: tf_eval_dataset[x] for x in tokenizer.model_input_names}
```

```
val_set_for_final_model = tf.data.Dataset.from_tensor_slices((eval_features, tf_eval_dataset["class"]))
```

```
val_set_for_final_model = val_set_for_final_model.batch(8)
```

```
test_features = {x: tf_test_dataset[x] for x in tokenizer.model_input_names}
```

```
test_set_for_final_model = tf.data.Dataset.from_tensor_slices((test_features, tf_test_dataset["class"]))
```

```
test_set_for_final_model = test_set_for_final_model.batch(8)
```

✓ Load and compile Model

```
model = TFAutoModelForSequenceClassification.from_pretrained("bert-base-cased", num_labels=3)
```

```
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=5e-5),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=tf.metrics.SparseCategoricalAccuracy(),
)
```

```
model.safetensors: 100% 436M/436M [00:02<00:00, 171MB/s]
All PyTorch model weights were used when initializing TFBertForSequenceClassification.
```

Some weights or buffers of the TF 2.0 model TFBertForSequenceClassification were not initialized from the PyTorch model and are newly initialized. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.


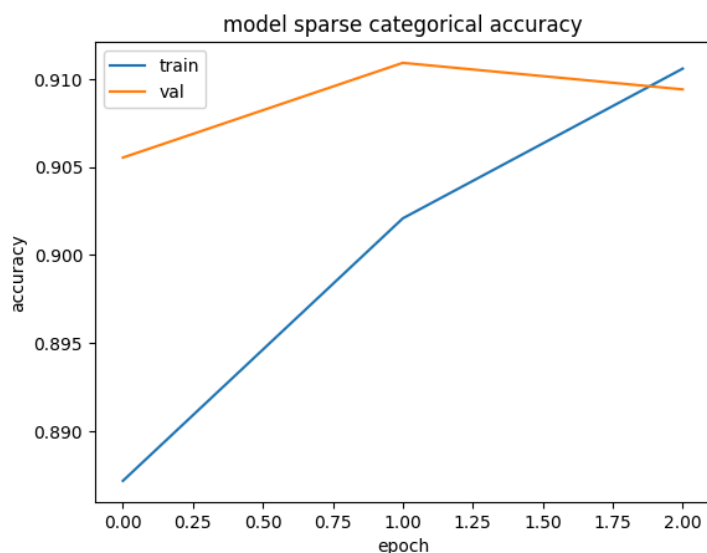
✓ Training Model

```
history = model.fit(train_set_for_final_model,
                    validation_data=val_set_for_final_model,
                    epochs=3 )
```

```
Epoch 1/3
2324/2324 [=====] - 2375s 987ms/step - loss: 0.3284 - sparse_categorical_accuracy: 0.8872 - val_loss: 0.2884 - val_sp
Epoch 2/3
2324/2324 [=====] - 2296s 988ms/step - loss: 0.2852 - sparse_categorical_accuracy: 0.9021 - val_loss: 0.2592 - val_sp
Epoch 3/3
2324/2324 [=====] - 2290s 985ms/step - loss: 0.2575 - sparse_categorical_accuracy: 0.9106 - val_loss: 0.2606 - val_sp
```

```
plt.plot(history.history['sparse_categorical_accuracy'])
plt.plot(history.history['val_sparse_categorical_accuracy'])
plt.title('model sparse categorical accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```



✓ Evaluate Model

```
test_loss, test_acc = model.evaluate(test_set_for_final_model, verbose=2)
print('\nTest accuracy:', test_acc)
```



194/194 - 63s - loss: 0.2516 - sparse_categorical_accuracy: 0.9116 - 63s/epoch - 325ms/step

Test accuracy: 0.911555826663971

✓ Predict Bert Model

```
predict_score_and_class_dict = {0: 'Hate Speech',
1: 'Offensive Language',
2: 'Neither'}
```

```
preds = model(tokenizer(["He is useless, I dont know why he came to our neighbourhood", "That guy sucks", "He is such a retard"], return_tensors="t"))
print(preds)
```

```
class_preds = np.argmax(preds, axis=1)
```

```
for pred in class_preds:
    print(predict_score_and_class_dict[pred])
```




```
tf.Tensor(
[[[-1.354622  0.49887672  0.855539 ]
 [-1.669037  0.07300647  1.7006899 ]
 [-0.24571076  2.3914983 -2.441549 ]], shape=(3, 3), dtype=float32)
Neither
Neither
Offensive Language
```

```
predict_score_and_class_dict = {0: 'Hate Speech',
1: 'Offensive Language',
```

```
2: 'Neither'}
preds = model(tokenizer(["He dresses up like a begger thise days"]),return_tensors="tf",padding=True,truncation=True))['logits']
print(preds)
class_preds = np.argmax(preds, axis=1)

for pred in class_preds:
    print(predict_score_and_class_dict[pred])
```



```
tf.Tensor([[ -1.5158917   0.4804067   1.0637728]], shape=(1, 3), dtype=float32)
Neither
```

Start coding or [generate](#) with AI.